

Mahalanobis OOD Detection for AI-Generated Text Classification

Dmitry Gorbunov

ITMO University, Institute of Applied Computer Science

Master's student in Artificial Intelligence, 1st year

2026

Abstract

We extend the AINL-Eval 2025 winning solution (sastsy, 91.22%) for detecting AI-generated scientific abstracts in Russian. The key challenge is identifying texts from unknown AI models not seen during training. We apply Mahalanobis distance-based OOD detection to Qwen2.5-7B with Dual-Head architecture, improving accuracy by +10.4% over softmax confidence (from 79.53% to 89.97%) with 76.25% unknown class recall. Mahalanobis also boosts lightweight ruBERT-tiny2 (29M params) to 85.25% – only 4.7% below Qwen, but 100x faster and 127x smaller. Code: <https://github.com/dpGorbunov/nlp-sem-project>.

1 Introduction

As LLMs become increasingly capable, distinguishing human-written scientific texts from AI-generated ones grows harder. LLMs can generate plausible abstracts that are difficult to distinguish from human-written text. This poses challenges for scientific publishing.

The AINL-Eval 2025 shared task [5] tackles this challenge for Russian scientific abstracts. The task is not just binary (human vs AI) – it requires identifying *which* AI model generated the text, including an "unknown" class for models not seen during training. This increases task complexity: the system must generalize to new AI models.

The winning solution by team sastsy achieved 91.22% accuracy using GigaCheck [1] with a Dual-Head modification. We extend this approach:

1. **Qwen2.5-7B**: stronger backbone than Mistral-7B on benchmarks
2. **Mahalanobis OOD Detection**: distance-based method for unknown class detection (+10.4% over baseline)
3. **Knowledge Distillation**: distillation to ruBERT-tiny2 for CPU inference

1.1 Team

Dmitry Gorbunov – model architecture design, experiments, report writing.

2 Related Work

GigaCheck and sastsy. The winning solution by team sastsy [5] built upon GigaCheck [1], which established a strong baseline using Mistral-7B fine-tuned with LoRA [3]. LoRA enables efficient adaptation by learning low-rank updates: $W' = W + \frac{\alpha}{r} \cdot BA$, where only matrices B and A are trained. Sastsy improved this approach with a dual-head architecture: binary classification (human vs AI) and multiclass classification (specific AI model). Weighted cross-entropy loss was used to handle class imbalance.

Hybrid approaches. Team adugeen [5] combined statistical and neural features. Their best test set results used bag-of-words features with binoculars [8] derived from Gemma-2B and LLaMA-1B. On the dev set, fine-tuning YandexGPT-8B with frozen backbone and trainable linear layers achieved the best performance.

3 Model Description

3.1 Architecture Overview

We follow the sastsy architecture but replace Mistral-7B with Qwen2.5-7B. We chose Qwen because, according to the Qwen2 Technical Report [2], it outperforms Mistral on standard benchmarks (Tab. 1), which may indicate better text understanding for this task.

The model has four components:

1. **Backbone:** Qwen2.5-7B fine-tuned with LoRA ($r=8$, $\alpha=16$) – we only train 0.04% of parameters
2. **Pooling:** EOS token embedding (the last token aggregates sequence information)
3. **Shared Layer:** Linear + tanh + dropout (transforms embeddings before classification)
4. **Dual-Head:** Two classification heads working together

Benchmark	Qwen2-7B	Mistral-7B	Δ
MMLU	70.3	64.2	+6.1
HumanEval	51.2	29.3	+21.9
GSM8K	79.9	52.2	+27.7

Table 1: Qwen2-7B vs Mistral-7B benchmark comparison.

3.2 Dual-Head Architecture

The key insight from `sastsy` is separating easy and hard tasks.

Binary Head answers: "Is this AI-generated?" This binary task is relatively easier – AI texts tend to exhibit detectable patterns.

Multiclass Head answers: "Which AI model?" The latter is more challenging – different LLMs produce similar outputs.

During training, both heads are optimized jointly:

$$\mathcal{L} = \mathcal{L}_{CE}^{bin} + \mathcal{L}_{CE}^{multi}$$

The multiclass loss ignores human samples (they have no AI model label).

During inference: if binary predicts "human" → output human. Otherwise, use multiclass prediction. The "unknown" class is not predicted directly – it is detected via Mahalanobis distance on embeddings.

3.3 Mahalanobis OOD Detection

The multiclass head can only predict known classes (GPT-4, Llama, Gemma). To detect unknown AI models, we use Mahalanobis distance [6] in embedding space:

$$D_M(x, c) = \sqrt{(x - \mu_c)^T \Sigma^{-1} (x - \mu_c)}$$

where Σ is the tied covariance matrix (shared across all classes). Unlike softmax confidence, which only considers output logits, it measures geometric distance to class centroids accounting for correlations. If a sample is far from all known classes, it is flagged as "unknown".

4 Dataset

The AINL-Eval 2025 dataset [5] contains Russian scientific abstracts from four sources: human-written, GPT-4-Turbo, Llama-3.3-70B, and Gemma-2-27B (Tab. 2).

	Train	Dev	Test
Samples	35,158	10,979	6,169
Classes	4	5	5

Table 2: AINL-Eval 2025 dataset statistics.

A key challenge: dev and test sets include a fifth class – "unknown" – generated by models *not present in training* (GigaChat-Lite in dev, DeepSeek-V3 in test). A classifier trained on four classes cannot directly recognize these models and must identify samples as AI-generated without assigning them to any known model.

This is the core OOD detection challenge. Standard softmax confidence performs poorly in this setting: the model confidently misclassifies unknown samples as one of the known AI models.

Notably, human texts are longer (126 words on average) and contain 10x more digits than AI-generated ones [5]. This suggests simple features could help, but our TF-IDF baseline shows they are not enough.

5 Experiments

5.1 Metrics

Primary metric: **Accuracy** (as per competition rules).

Detailed per-class precision, recall, and F1-score are available in the accompanying notebook; here we focus on accuracy and visualize results with confusion matrices.

5.2 Experiment Setup

- GPU: NVIDIA A100 40GB
- Precision: bfloat16
- Batch: 16
- Learning rate: 3e-5
- Epochs: 10, Early stopping: patience=3
- LoRA: r=8, alpha=16, targets: q_proj, v_proj
- Mahalanobis threshold: $\tau=77.5$, selected via grid search on dev set

5.3 Baselines

- TF-IDF + Logistic Regression
- ruBERT-tiny2 fine-tuned
- sastsy [5]: 1st place winner (GigaCheck-based)

6 Results

Tab. 3 shows our main results. We test three OOD detection strategies: no detection (Base), softmax confidence threshold, and Mahalanobis distance. We also experimented with Energy score and Combined (Energy + Confidence), but they performed worse (Fig. 1).

Mahalanobis outperforms confidence-based methods because a model trained only on GPT-4/Llama/Gemma will confidently assign GigaChat samples to one of these classes – it has no concept of "none of the above." Mahalanobis detects unknown AI models because they produce representations far from all training classes.

Key findings:

Method	Base	Confidence	Mahalanobis
<i>AINL-Eval 2025 results [5]:</i>			
sastsy (1st place)	91.22%	—	—
adugeen (2nd place)	86.96%	—	—
TF-IDF baseline	80.81%	—	—
<i>Our experiments:</i>			
Qwen2.5 + Dual-Head (ours)	79.53%	82.61%	89.97%
ruBERT-tiny2 (fine-tuned)	78.29%	80.29%	85.25%
TF-IDF + LogReg	76.85%	81.06%	—

Table 3: Comparison of methods with different OOD detection strategies. All results evaluated on dev set.

1. This approach boosts accuracy from 79.53% to 89.97% (+10.4%). The binary head alone achieves 95.38% binary accuracy (5 classes collapsed to human/AI), and unknown recall reaches 76.25%.
2. ruBERT-tiny2 achieves 85.25% – only 4.7% below Qwen (see Table 4 for details). This makes real-time CPU deployment feasible.




Figure 1: Mahalanobis OOD detection results. Left: confusion matrix. Center: distance distribution by class with threshold $\tau=77.5$. Right: accuracy comparison.

6.1 Knowledge Distillation

Qwen2.5-7B requires GPU and is too slow for real-time applications. Following DisRanker [4], which showed LLM knowledge can be distilled to BERT with 10x speedup, we compress to ruBERT-tiny2 (29M params):

$$\mathcal{L} = \alpha \cdot \mathcal{L}_{KL}(p_s, p_t) \cdot T^2 + (1 - \alpha) \cdot \mathcal{L}_{CE}(p_s, y)$$

where $T = 4$ (temperature), $\alpha = 0.7$.

We compare two approaches: distillation without prior fine-tuning (ruBERT-tiny2 + KD) and distillation after fine-tuning (ruBERT-tiny2 FT + KD). Results are shown in Table 4.

Model	Size	Inference	Raw Acc	+Mahalanobis
Qwen2.5-7B (teacher)	15 GB	~300ms (GPU)	79.53%	89.97%
ruBERT-tiny2 (fine-tuned)	118 MB	~3ms (GPU)	78.29%	85.25%
ruBERT-tiny2 + KD	118 MB	~3ms (GPU)	74.21%	80.44%
ruBERT-tiny2 FT + KD	118 MB	~3ms (GPU)	77.01%	85.25%

Table 4: Teacher vs Student comparison. FT = fine-tuned, KD = Knowledge Distillation. Qwen: 7.61B params [2], ruBERT-tiny2: 29M params [7] (260× fewer params, 127× smaller file size). Inference times based on [7].

Observation: Distillation without prior fine-tuning achieves lower accuracy than the fine-tuning baseline (74.21% vs 78.29%). This is expected – the student model needs more capacity or training to learn from scratch. Applying KD after fine-tuning (ruBERT-tiny2 FT + KD) achieves the same final accuracy as the baseline (85.25% with Mahalanobis), suggesting the model has already converged and KD provides no additional benefit in this setting.

7 Conclusion

The main contribution of this work: **detecting unknown AI models is challenging; Mahalanobis distance provides an effective solution.** Standard confidence-based methods perform poorly because neural networks tend to be overconfident on out-of-distribution samples. Mahalanobis offers a more robust criterion: "unknown" means "far from everything I've seen".

Our best model (Qwen2.5-7B + Dual-Head + Mahalanobis) achieves 89.97% accuracy, with 76.25% recall on the unknown class. Notably, ruBERT-tiny2 – a model 260x smaller – achieves comparable results with the same approach. This suggests that the OOD detection method may matter more than model size for this task.

Practical implications: A 29M parameter model running in 3ms on GPU (or 5ms on CPU) can detect AI-generated scientific abstracts with reasonable accuracy. This enables deployment in resource-constrained environments - browser extensions, email filters, or mobile apps - without GPU infrastructure.

Limitations: The method requires computing class statistics from training data embeddings upfront. If the distribution of AI models shifts (new models appear), these statistics need recomputation. Future work could explore online or adaptive OOD detection methods.

References

- [1] Tolstykh, I., Tsybina, A., Yakubson, S., Gordeev, A., Dokholyan, V., Kuprashevich, M. (2024). GigaCheck: Detecting LLM-generated Content. *arXiv preprint arXiv:2410.23728*.

- [2] Qwen Team. (2024). Qwen2 Technical Report. *arXiv preprint arXiv:2407.10671*.
- [3] Hu, E. J., et al. (2022). LoRA: Low-Rank Adaptation of Large Language Models. *ICLR 2022*.
- [4] Ye, D., et al. (2024). Best Practices for Distilling Large Language Models into BERT for Web Search Ranking. *arXiv preprint arXiv:2411.04539*.
- [5] Batura, T., Bruches, E., Shvenk, M., Malykh, V. (2025). AINL-Eval 2025 Shared Task: Detection of AI-Generated Scientific Abstracts in Russian. *arXiv preprint arXiv:2508.09622*. <https://codalab.lisn.upsaclay.fr/competitions/21895>
- [6] Lee, K., et al. (2018). A Simple Unified Framework for Detecting Out-of-Distribution Samples and Adversarial Attacks. *NeurIPS 2018*.
- [7] Dale, D. (2022). ruBERT-tiny2: Russian Sentence Encoder. *Habr*. <https://habr.com/ru/post/669674/>, <https://huggingface.co/cointegrated/rubert-tiny2>
- [8] Hans, A., et al. (2024). Spotting LLMs With Binoculars: Zero-Shot Detection of Machine-Generated Text. *ICML 2024*. <https://arxiv.org/abs/2401.12070>